

Name:Nitin Panwar CU240251575 sec:B

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder, StandardScaler

print("libraries are installed ") libraries are installed

df =
pd.read_csv("/workspaces/Mastery_in_DataAnalytics_and_Visualizations_and_Career_Advancemen/train-selected-columns.csv")

print("First 5 rows:")
print(df.head())

print("\nDataset Information:")
df.info()

print("\nStatistical Description:")
print(df.describe())
```

First 5 rows:

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

		Name	Sex	Age	
SibSp	\				
		Braund, Mr. Owen Harris	male	22.0	0
		Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1
		Heikkinen, Miss. Laina	female	26.0	2
		Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	3
		Allen, Mr. William Henry	male	35.0	4
Parch					
0	0	A/5 21171		7.2500	
1	0	PC 17599		71.2833	

2	0	STON/O2. 3101282		7.9250	
---	---	------------------	--	--------	--

```
3      0      113803  53.1000
4      0      373450   8.0500
```

Dataset Information:

```
<class 'pandas.DataFrame'>
```

RangeIndex: 891 entries, 0 to 890

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	str
4	Sex	891 non-null	str
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64
7	Parch	891 non-null	int64
8	Ticket	891 non-null	str
9	Fare	891 non-null	float64

dtypes: float64(2), int64(5), str(3) memory usage: 69.7 KB

Statistical Description:

	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	714.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008
std	257.353842	0.486592	0.836071	14.526497	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
print("Dataset Information:")
df.info()
```

Dataset Information:

```
<class 'pandas.DataFrame'>
```

RangeIndex: 891 entries, 0 to 890

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	str
4	Sex	891 non-null	str
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64
7	Parch	891 non-null	int64
8	Ticket	891 non-null	str
9	Fare	891 non-null	float64

float64 dtypes: float64(2), int64(5), str(3) memory usage: 69.7 KB

```
print("Statistical Description:")
```

```
df.describe()
```

Statistical

Description:

	PassengerId	Survived	Pclass
Age	SibSp	\	
count	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642
std	257.353842	29.699118	0.523008
min	1.000000	0.000000	1.000000
25%	0.000000	0.420000	0.000000
50%	446.000000	0.000000	20.125000
75%	668.500000	0.000000	0.000000
max	891.000000	1.000000	3.000000

```
Parch      Fare
count
891.000000
891.000000  mean
0.381594
32.204208   std
0.806057
49.693429   min
0.000000
0.000000    25%
0.000000
7.910400
50%      0.000000
14.454200  75%
0.000000
31.000000   max
6.000000
512.329200
print("Dataset Shape:")
print(df.shape)
Dataset Shape:
(891, 10)

print("Missing Values:")
print(df.isnull().sum())
```

Missing Values:

PassengerId	0
Survived	0
Pclass	0
Name	0
Sex	0
Age	177
SibSp	0
Parch	0
Ticket	0
Fare	0

dtype: int64

```
print("Missing Value Percentage:")
```

```
print((df.isnull().sum() / len(df)) * 100)
```

Missing Value Percentage:

PassengerId	0.00000
Survived	0.00000
Pclass	0.00000
Name	0.00000
Sex	0.00000
Age	19.86532
SibSp	0.00000
Parch	0.00000
Ticket	0.00000
Fare	0.00000

dtype: float64

```
print("Duplicate records before removal:")
```

```
print(df.duplicated().sum())
```

```
df.drop_duplicates(inplace=True)
```

```
print("Duplicate records after removal:")
```

```
print(df.duplicated().sum())
```

Duplicate records before removal:

0 Duplicate records after
removal:
0

df =

```
pd.read_csv("/workspaces/Mastery_in_DataAnalytics_and_Visualizations_and_Career_Advancemen/train-selected-columns.csv")
```

```
print(df.columns.tolist())
```

```
['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch', 'Ticket', 'Fare']
```

```

df.to_csv("cleaned_train.csv", index=False)
print("Cleaned dataset saved successfully!")
Cleaned dataset saved successfully!
# Check duplicate records
print("Duplicate records before removal:")
print(df.duplicated().sum())

# Remove duplicate records
df.drop_duplicates(inplace=True)

# Check duplicates again
print("\nDuplicate records after removal:")
print(df.duplicated().sum())

Duplicate records before removal:
0

Duplicate records after removal:
0 # Step 7: Identify outliers using Box Plot and
IQR

import matplotlib.pyplot as plt
import seaborn as sns

# Box plots
plt.figure(figsize=(10, 5))
sns.boxplot(data=df[['Age', 'Fare']])
plt.title("Box Plot for Age and Fare")
plt.show()

# IQR for Age
Q1 = df['Age'].quantile(0.25)
Q3 = df['Age'].quantile(0.75)

IQR = Q3 - Q1


lower_limit = Q1 - 1.5 * IQR
upper_limit = Q3 + 1.5 * IQR

# Remove Age outliers
df = df[(df['Age'].isna()) |
        ((df['Age'] >= lower_limit) & (df['Age'] <= upper_limit))]

print("Outliers in Age treated successfully.")
print("New dataset shape:", df.shape)

```

Box Plot for Age and Fare



Outliers in Age treated successfully.
New dataset shape: (880, 10)

Step 8: Normalization

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
```

Age and Fare normalized successfully.

```
df[['Age', 'Fare']] = scaler.fit_transform(df[['Age', 'Fare']])
```

```
print("Age and Fare normalized successfully.")
```

```
print(df[['Age', 'Fare']].head())
```

```
4    0.543882    0.015713
```

Step 9: Feature Engineering

```
0    0.339415    0.014151 # Create
```

```
1    0.591066    0.139136 Family
```

```
2    0.402328    0.015469
```

```
3    0.543882    0.103644
```

Size

```
df['FamilySize'] = df['SibSp'] + df['Parch'] + 1
```

Create Age Group

```
df['AgeGroup'] = pd.cut(
```

```
df['Age'],
```

```

    bins=[0, 18, 35, 60, 100],
    labels=['Child', 'Young', 'Adult', 'Senior']
)

print("New features created successfully.")
print(df[['FamilySize', 'AgeGroup']].head())
New features created successfully.
   FamilySize AgeGroup
0           2    Child
1           2    Child
2           1    Child
3           2    Child
4           1    Child
# Step 10: Save the cleaned dataset

df.to_csv('titanic_cleaned.csv', index=False)

print("Cleaned Titanic dataset saved successfully!")

Cleaned Titanic dataset saved successfully!

```